

0. Foreword

This purpose of this page is to provide documentation for the DVI file format. Much of the content on this page was taken directly from the `dvitype.web` program, written by Donald Knuth (this explains the several references to the DVItype program; these can be ignored).

1. The DVI File Format

Before we get into the details of DVItype, we need to know exactly what DVI files are. The form of such files was designed by David R. Fuchs in 1979. Almost any reasonable typesetting device can be driven by a program that takes DVI files as input, and dozens of such DVI-to-whatever programs have been written. Thus, it is possible to print the output of document compilers like TeX on many different kinds of equipment.

A DVI file is a stream of 8-bit bytes, which may be regarded as a series of commands in a machine-like language. The first byte of each command is the operation code, and this code is followed by zero or more bytes that provide parameters to the command. The parameters themselves may consist of several consecutive bytes; for example, the `set_rule` command has two parameters, each of which is four bytes long. Parameters are usually regarded as nonnegative integers; but four-byte-long parameters, and shorter parameters that denote distances, can be either positive or negative. Such parameters are given in two's complement notation. For example, a two-byte-long distance parameter has a value between -2^{15} and $2^{15}-1$. **[NOTE: DVI files use big endian format for multiple byte integer parameters.]**

A DVI file consists of a ```preamble,` followed by a sequence of one or more ```pages,` followed by a ```postamble.` The preamble is simply a `pre` command, with its parameters that define the dimensions used in the file; this must come first. Each ```page` consists of a `bop` command, followed by any number of other commands that tell where characters are to be placed on a physical page, followed by an `eop` command. The pages appear in the order that they were generated, not in any particular numerical order. If we ignore `nop` commands and `fnt_def` commands (which are allowed between any two commands in the file), each `eop` command is immediately followed by a `bop` command, or by a `post` command; in the latter case, there are no more pages in the file, and the remaining bytes form the postamble. Further details about the postamble will be explained later.

Some parameters in DVI commands are ```pointers.` These are four-byte quantities that give the location number of some other byte in the file; the first byte is number 0, then comes number 1, and so on. For example, one of the parameters of a `bop` command points to the previous `bop`; this makes it feasible to read the pages in backwards order, in case the results are being directed to a device that stacks its output face up. Suppose the preamble of a DVI file occupies bytes 0 to 99. Now if the first page occupies bytes 100 to 999, say, and if the second page occupies bytes 1000 to 1999, then the `bop` that starts in byte 1000 points to 100 and the `bop` that starts in byte 2000 points to 1000. (The very first `bop`, i.e., the one that starts in byte 100, has a pointer of -1.)

The DVI format is intended to be both compact and easily interpreted by a machine. Compactness is achieved by making most of the information implicit instead of explicit. When a DVI-reading program reads the commands for a page, it keeps track of several quantities: (a) The current font `f` is an integer; this value is changed only by `fnt` and `fnt_num` commands. (b) The current position on the page is given by two numbers called the horizontal and vertical coordinates, `h` and `v`. Both coordinates are zero at the upper left corner of the page; moving to the right corresponds to increasing the horizontal coordinate, and moving down corresponds to increasing the vertical coordinate. Thus, the coordinates are essentially Cartesian, except that vertical directions are flipped; the Cartesian version of `(h, v)` would be `(h, -v)`. (c) The current spacing amounts are given by four numbers `w`, `x`, `y`,

The values of h , v , w , x , y , and z are signed integers having up to 32 bits, including the sign. Since they represent physical distances, there is a small unit of measurement such that increasing h by 1 means moving a certain tiny distance to the right. The actual unit of measurement is variable, as explained below.

3. Table of Opcodes

The following table gives the instruction set for DVI. The parameters are listed in the order they would appear in a DVI file; the number in brackets gives the size of the parameter (in bytes).

The DVI Instruction Set			
Opcode	Instruction Name	Parameters	Description
0...127	set_char_i		typeset a character and move right
128	set1	c[1]	typeset a character and move right
129	set2	c[2]	
130	set3	c[3]	
131	set4	c[4]	
132	set_rule	a[4], b[4]	typeset a rule and move right
133	put1	c[1]	typeset a character
134	put2	c[2]	
135	put3	c[3]	
136	put4	c[4]	
137	put_rule	a[4], b[4]	typeset a rule
138	nop		no operation
139	bop	c_0[4]...c_9[4], p[4]	beginning of page
140	eop		ending of page
141	push		save the current positions
142	pop		restore previous positions
143	right1	b[1]	move right
144	right2	b[2]	
145	right3	b[3]	
146	right4	b[4]	
147	w0		move right by w
148	w1	b[1]	move right and set w
149	w2	b[2]	
150	w3	b[3]	
151	w4	b[4]	
152	x0		move right by x
153	x1	b[1]	move right and set x

157	down1	a[1]	move down
158	down2	a[2]	
159	down3	a[3]	
160	down4	a[4]	
161	y0		move down by y
162	y1	a[1]	move down and set y
163	y2	a[2]	
164	y3	a[3]	
165	y4	a[4]	
166	z0		move down by z
167	z1	a[1]	move down and set z
168	z2	a[2]	
169	z3	a[3]	
170	z4	a[4]	
171...234	fnt_num_i		set current font to i
235	fnt1	k[1]	set current font
236	fnt2	k[2]	
237	fnt3	k[3]	
238	fnt4	k[4]	
239	xxx1	k[1], x[k]	extension to DVI primitives
240	xxx2	k[2], x[k]	
241	xxx3	k[3], x[k]	
242	xxx4	k[4], x[k]	
243	fnt_def1	k[1], c[4], s[4], d[4], a[1], l[1], n[a+l]	define the meaning of a font number
244	fnt_def2	k[2], c[4], s[4], d[4], a[1], l[1], n[a+l]	
245	fnt_def3	k[3], c[4], s[4], d[4], a[1], l[1], n[a+l]	
246	fnt_def4	k[4], c[4], s[4], d[4], a[1], l[1], n[a+l]	
247	pre	i[1], num[4], den[4], mag[4], k[1], x[k]	preamble
248	post	p[4], num[4], den[4], mag[4], l[4], u[4], s[2], t[2] < font definitions >	postamble beginning
249	post_post	q[4], i[1]; 223's	postamble ending
250...255	<i>undefined</i>		

4. Description of Opcodes

cannot be sure that h will advance after this command; but h usually does increase.

- Opcodes 128-131: `seti (1 ≤ i ≤ 4); c[i]`

Same as `set_char_0`, except that character number c is typeset. TeX82 uses the `set1` command for characters in the range $128 ≤ c < 256$. TeX82 never uses the `set2`, command which is intended for processors that deal with oriental languages; but DVItypset will allow character codes greater than 255, assuming that they all have the same width as the character whose code is $c \bmod 256$.

- Opcode 132: `set_rule; a[4], b[4]`

Typeset a solid black rectangle of height a and width b , with its bottom left corner at (h, v) . Then set $h := h + b$. If either $a ≤ 0$ or $b ≤ 0$, nothing should be typeset. Note that if $b < 0$, the value of h will decrease even though nothing else happens. Programs that typeset from DVI files should be careful to make the rules line up carefully with digitized characters, as explained in connection with the `rule_pixels` subroutine below.

- Opcodes 133-136: `puti (1 ≤ i ≤ 4); c[i]`

Typeset character number c from font f such that the reference point of the character is at (h, v) . (The `put` commands are exactly like the `set` commands, except that they simply put out a character or a rule without moving the reference point afterwards.)

- Opcode 137: `put_rule; a[4], b[4]`

Same as `set_rule`, except that h is not changed.

- Opcode 138: `nop`

No operation, do nothing. Any number of `nop`'s may occur between DVI commands, but a `nop` cannot be inserted between a command and its parameters or between two parameters.

- Opcode 139: `bop; c_0[4]..c_9[4], p[4]`

Beginning of a page: Set $(h, v, w, x, y, z) := (0, 0, 0, 0, 0, 0)$ and set the stack empty. Set the current font f to an undefined value. The ten c_i parameters can be used to identify pages, if a user wants to print only part of a DVI file; TeX82 gives them the values of `\count0... \count9` at the time `\shipout` was invoked for this page. The parameter p points to the previous `bop` command in the file, where the first `bop` has $p = -1$.

- Opcode 140: `eop`

End of page: Print what you have read since the previous `bop`. At this point the stack should be empty. (The DVI-reading programs that drive most output devices will have kept a buffer of the material that appears on the page that has just ended. This material is largely, but not entirely, in order by v coordinate and (for fixed v) by h coordinate; so it usually needs to be sorted into some order that is appropriate for the device in question. DVItypset does not do such sorting.)

- Opcode 141: `push`

Push the current values of (h, v, w, x, y, z) onto the top of the stack; do not change any of these values. Note that f is not pushed.

- Opcode 142: `pop`

Pop the top six values off of the stack and assign them to (h, v, w, x, y, z) . The number of pops should never exceed the number of pushes, since it would be highly embarrassing if

Set $h:=h+b$, i.e., move right b units. The parameter is a signed number in two's complement notation; if $b < 0$, the reference point actually moves left.

- Opcodes 147-151: wi ($0 \leq i \leq 4$); $b[i]$

The $w0$ command sets $h:=h+w$; i.e., moves right w units. With luck, this parameterless command will usually suffice, because the same kind of motion will occur several times in succession. The other w commands set $w:=b$ and $h:=h+b$. The value of b is a signed quantity in two's complement notation. This command changes the current w spacing and moves right by b .

- Opcodes 152-156: xi ($0 \leq i \leq 4$); $b[i]$

The parameterless $x0$ command sets $h:=h+x$; i.e., moves right x units. The x commands are like the w commands except that they involve x instead of w . The other x commands set $x:=b$ and $h:=h+b$. The value of b is a signed quantity in two's complement notation. This command changes the current x spacing and moves right by b .

- Opcodes 157-160: $downi$ ($1 \leq i \leq 4$); $a[i]$

Set $v:=v+a$, i.e., move down a units. The parameter is a signed number in two's complement notation; if $a < 0$, the reference point actually moves up.

- Opcodes 161-165: yi ($0 \leq i \leq 4$); $a[i]$

The $y0$ command sets $v:=v+y$; i.e., moves down y units. With luck, this parameterless command will usually suffice, because the same kind of motion will occur several times in succession. The other y commands set $y:=a$ and $v:=v+a$. The value of a is a signed quantity in two's complement notation. This command changes the current y spacing and moves down by a .

- Opcodes 166-170: zi ($0 \leq i \leq 4$); $a[i]$

The $z0$ command sets $v:=v+z$; i.e., moves down z units. The z commands are like the y commands except that they involve z instead of y . The other z commands set $z:=a$ and $v:=v+a$. The value of a is a signed quantity in two's complement notation. This command changes the current z spacing and moves down by a .

- Opcodes 171-234: fnt_num_i ($0 \leq i \leq 63$)

Set $f:=i$. Font i must previously have been defined by a fnt_def instruction, as explained below.

- Opcodes 235-238: $fnti$ ($1 \leq i \leq 4$); $k[i]$

Set $f:=k$. TeX82 uses the $fnt1$ command for font numbers in the range $64 \leq k < 256$. TeX82 never generates the $fnt2$ command, but large font numbers may prove useful for specifications of color or texture, or they may be used for special fonts that have fixed numbers in some external coding scheme.

- Opcodes 239-242: $xxxi$ ($1 \leq i \leq 4$); $k[i]$, $x[k]$

This command is undefined in general; it functions as a $k+i+1$ -byte nop unless special DVI-reading programs are being used. TeX82 generates $xxx1$ when a short enough `\special` appears, setting k to the number of bytes being sent. It is recommended that x be a string having the form of a keyword followed by possible parameters relevant to that keyword.

found by programs that read this DVI file.

Parameter *s* contains a fixed-point scale factor that is applied to the character widths in font *k*; font dimensions in TFM files and other font files are relative to this quantity, which is always positive and less than 2^{27} . It is given in the same units as the other dimensions of the DVI file. Parameter *d* is similar to *s*; it is the ``design size,' ' and (like *s*) it is given in DVI units. Thus, font *k* is to be used at $\text{mag } s / 1000 \text{ } d$ times its normal size.

The remaining part of a font definition gives the external name of the font, which is an ASCII string of length *a+l*. The number *a* is the length of the ``area' ' or directory, and *l* is the length of the font name itself; the standard local system font area is supposed to be used when *a*=0. The *n* field contains the area in its first *a* bytes.

Font definitions must appear before the first use of a particular font number. Once font *k* is defined, it must not be defined again; however, we shall see below that font definitions appear in the postamble as well as in the pages, so in this sense each font number is defined exactly twice, if at all. Like *nop* commands, font definitions can appear before the first *bop*, or between an *eop* and a *bop*.

- Opcodes 247: *pre*; *i*[1], *num*[4], *den*[4], *mag*[4], *k*[1], *x*[*k*]

The preamble contains basic information about the file as a whole and must come at the very beginning of the file. The *i* byte identifies DVI format; currently this byte is always set to 2. (The value *i*=3 is currently used for an extended format that allows a mixture of right-to-left and left-to-right typesetting. Some day we will set *i*=4, when DVI format makes another incompatible change - perhaps in the year 2048.)

The next two parameters, *num* and *den*, are positive integers that define the units of measurement; they are the numerator and denominator of a fraction by which all dimensions in the DVI file could be multiplied in order to get lengths in units of $10^{(-7)}$ meters. (For example, there are exactly 7227 TeX points in 254 centimeters, and TeX82 works with scaled points where there are 2^{16} sp in a point, so TeX82 sets *num*=25400000 and *den*=7227 2^{16} =473628672.

The *mag* parameter is what TeX82 calls `\mag`, i.e., 1000 times the desired magnification. The actual fraction by which dimensions are multiplied is therefore $m \text{ } n / 1000 \text{ } d$. Note that if a TeX source document does not call for any true dimensions, and if you change it only by specifying a different `\mag` setting, the DVI file that TeX creates will be completely unchanged except for the value of *mag* in the preamble and postamble. (Fancy DVI-reading programs allow users to override the *mag* setting when a DVI file is being printed.)

Finally, *k* and *x* allow the DVI writer to include a comment, which is not interpreted further. The length of comment *x* is *k*, where $0 \leq k < 256$.

- Opcodes 248: *post*; *p*[4], *num*[4], *den*[4], *mag*[4], *l*[4], *u*[4], *s*[2], *t*[2]; < *font definitions* >

The last page in a DVI file is followed by *post*; this command introduces the postamble, which summarizes important facts that TeX has accumulated about the file, making it possible to print subsets of the data with reasonable efficiency. The parameter *p* is a pointer to the final *bop* in the file. The next three parameters, *num*, *den*, and *mag*, are duplicates of the quantities that appeared in the preamble.

Parameters *l* and *u* give respectively the height-plus-depth of the tallest page and the width of the widest page, in the same units as other dimensions of the file. These numbers might be used by a DVI-reading program to position individual ``pages' ' on large sheets of film or paper; however, the standard convention for output on normal size paper is to position each page so that the upper left-hand corner is exactly one inch from the left and the top. Experience has shown that it is unwise to design DVI-to-printer software that attempts cleverly to center the output; a fixed position of

~~pop commands) needed to process this file. Then comes 4, the total number of pages (pop commands) present.~~

The postamble continues with font definitions, which are any number of `fnt_def` commands as described above, possibly interspersed with `nop` commands. Each font number that is used in the DVI file must be defined exactly twice: Once before it is first selected by a `fnt` command, and once in the postamble.

- Opcodes 249: `post_post; q[4], i[1]; 223's`

The last part of the postamble, following the `post_post` byte that signifies the end of the font definitions, contains `q` a pointer to the `post` command that started the postamble. An identification byte, `i`, comes next; this currently equals 2, as in the preamble.

The `i` byte is followed by four or more bytes that are all equal to the decimal number 223 (i.e., 337 in octal). TeX puts out four to seven of these trailing bytes, until the total length of the file is a multiple of four bytes, since this works out best on machines that pack four bytes per word; but any number of 223's is allowed, as long as there are at least four of them. In effect, 223 is a sort of signature that is added at the very end.

This curious way to finish off a DVI file makes it feasible for DVI-reading programs to find the postamble first, on most computers, even though TeX wants to write the postamble last. Most operating systems permit random access to individual words or bytes of a file, so the DVI reader can start at the end and skip backwards over the 223's until finding the identification byte. Then it can back up four bytes, read `q`, and move to byte `q` of the file. This byte should, of course, contain the value 248 (`post`); now the postamble can be read, so the DVI reader discovers all the information needed for typesetting the pages. Note that it is also possible to skip through the DVI file at reasonably high speed to locate a particular page, if that proves desirable. This saves a lot of time, since DVI files used in production jobs tend to be large.